

HW5 Responses

Yawei LI

2/27/2020

Question 1

When growing a classification tree, Gini index and the cross-entropy are expected to do a better job than plain error rate because they react faster to impurity in classification.

When pruning a decision tree, classification error will come in handy because they are closer to the very concept of “correct classification” that we try to achieve.

Question 2

Data and Prep

```
train = read.csv('gss_train.csv')
test = read.csv('gss_test.csv')
train$colrac = as.factor(train$colrac)
test$colrac = as.factor(test$colrac)
models = c("Logistic regression",
           'Naive Bayes',
           'Elastic net regression',
           'Decision tree (CART)',
           'Bagging',
           'Random forest',
           'Boosting')
cv_error = numeric()
auc = numeric()
```

Logistic

```
set.seed(1234)
train.control <- trainControl(method = "cv", number = 10)
logit <- train(colrac ~., data = train, method = "glm", family = 'binomial',
               trControl = train.control)
cv_error[1] <- logit$results$Accuracy
logit_pred <- as.numeric(predict(logit, newdata = train))
auc[1] <- roc(train$colrac, logit_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

naive Bayes

```
set.seed(2234)
train.control <- trainControl(method = "cv", number = 10)
nb <- train(colrac ~., data = train, method = "nb",
            trControl = train.control)
cv_error[2] <- nb$results$Accuracy
#I got 2 accuracies here labeled for T and F, not sure which to use.
nb_pred <- as.numeric(predict(nb,newdata = train))
auc[2] <- roc(train$colrac, nb_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

Elastic Net

```
set.seed(3234)
train.control <- trainControl(method = "cv", number = 10)
enet <- train(colrac ~., data = train, method = "glmnet",
             trControl = train.control)
cv_error[3] <- enet$results$Accuracy[5] # Chosen by reading summary
enet_pred <- as.numeric(predict(enet,newdata = train))
auc[3] <- roc(train$colrac, enet_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

CART

```
set.seed(4234)
cart <- tree(colrac ~., data = train)
cart_cv <- cv.tree(cart, FUN = prune.misclass)
min_idx <- which.min(cart_cv$dev)
best_size <- cart_cv$size[min_idx]
cart_prune <- prune.misclass(cart, best = best_size)
cv_error[4] <- 1 - summary(cart_prune)$misclass[1]/summary(cart_prune)$misclass[2]
# Misclassifications / Total number of obs.
cart_pred <- as.numeric(predict(cart_prune,newdata = train, type = 'class'))
auc[4] <- roc(train$colrac, cart_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

Bagging

```
set.seed(5234)
ctrl <- trainControl(method = "cv", number = 10)
bagged <- train(
  colrac ~ .,
```

```

data = train,
method = "treebag",
trControl = ctrl
)
cv_error[5] <- bagged$results$Accuracy
bagged_pred <- as.numeric(predict(bagged,newdata = train))
auc[5] <- roc(train$colrac, bagged_pred)$auc[1]

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

```

Random Forest

```

set.seed(6234)
rf1 <- train(
  colrac ~ .,
  data = train,
  method = "ranger",
  trControl = ctrl,
  # splitrule = 'gini' # This is also causing error same as below.
  importance = "impurity"
) # Maxi Accuracy 0.801 at 3rd model

# I am tring to tuning node size here.
# But encountered an error of
# formal argument min.node.size matched by multiple actual arguments
# by adding min.node.size = 3, default is 1.
# The caret documentation said it could tune min.node.size, yet
# I got it set at default.

#set.seed(7234)
#rf1 <- train(
# colrac ~ .,
# data = train,
# method = "ranger",
# trControl = ctrl,
# importance = "impurity",
# min.node.size = 3
#)

cv_error[6] <- rf1$results$Accuracy[3]
rf_pred <- as.numeric(predict(rf1,newdata = train))
auc[6] <- roc(train$colrac, rf_pred)$auc[1]

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

```

Boosting

```
set.seed(7234)
boost <- train(
  colrac ~ .,
  data = train,
  method = "gbm",
  trControl = trainControl(method = "cv",
    number = 10,
    verboseIter = FALSE),
  verbose = 0)
cv_error[7] <- boost$results$Accuracy[8] # Max accuracy 0.809 at 8th model
boost_pred <- as.numeric(predict(boost, newdata = train))
auc[7] <- roc(train$colrac, boost_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

Question 3 & 4 Model Comparison

As can be seen from the table below, Boosting model provides the best cross-validation correct classification rate, while maintains a level of area under curve (the best if we ignore the seemingly problematic bagging and random forest auc). Therefore, I believe boosting model performs the best according to our rubrics.

```
report <- data.frame(models, cv_error, auc)
print(report)
```


##	models	cv_error	auc
## 1	Logistic regression	0.7994346	0.8164760
## 2	Naive Bayes	0.7371760	0.7431246
## 3	Elastic net regression	0.8021925	0.8179706
## 4	Decision tree (CART)	0.7621951	0.7535254
## 5	Bagging	0.7784657	0.9972150
## 6	Random forest	0.8001178	1.0000000
## 7	Boosting	0.8089125	0.8505333

As can be seen from the feature importance plot and table below. A 44 feature model may seem over complicated to a certain point, yet a closer examination reveals that the model does drop quite a number of insignificant features and assigned very different weights to important and unimportant features. Therefore, though it may suffer from a certain degree of over-fitting, it also takes better advantage of the large number of features in our gss data, and limits the effects of over fitting to a smaller degree.

```
print(boost$finalModel)
```

```
## A gradient boosted model with bernoulli loss function.
## 150 iterations were performed.
## There were 55 predictors of which 44 had non-zero influence.
```

```
print(summary(boost$finalModel))
```



##	var	rel.inf
## tolerance	tolerance	43.46272149
## colmslm	colmslm	9.71902432
## colath	colath	8.68012869
## colmil	colmil	5.49602033
## age	age	4.60528792
## sibs	sibs	2.12984569
## income06	income06	2.00327861
## science_quiz	science_quiz	1.97087363
## social_connect	social_connect	1.47990249
## egalit_scale	egalit_scale	1.41783459
## polviews	polviews	1.36620771
## wordsum	wordsum	1.30502040
## con_govt	con_govt	1.01723045
## authoritarianism	authoritarianism	1.00017260
## tvhours	tvhours	0.94219244
## mode	mode	0.92940211
## colcom	colcom	0.88987921
## homosex	homosex	0.84722764
## marital_Never.married	marital_Never.married	0.78418141
## pray	pray	0.76484468
## south	south	0.72516587
## social_cons3_Conserv	social_cons3_Conserv	0.69252993
## childs	childs	0.62554991
## grass	grass	0.60915174
## colhomo	colhomo	0.52800977

```
## happy                happy 0.52345096
## attend               attend 0.52250550
## evangelical          evangelical 0.50387497
## owngun              owngun 0.46616147
## marital_other        marital_other 0.44275076
## degree_other         degree_other 0.41860186
## partyid_3_Ind        partyid_3_Ind 0.39904746
## spend3_Mod           spend3_Mod 0.35734474
## teensex             teensex 0.35639610
## vetyears            vetyears 0.34332031
## partyid_3_Rep        partyid_3_Rep 0.33451337
## marital_Divorced     marital_Divorced 0.33171433
## reborn_r            reborn_r 0.23580795
## relig_other          relig_other 0.17463947
## spend3_Liberal       spend3_Liberal 0.15264655
## hispanic_2          hispanic_2 0.14093224
## pres08              pres08 0.11064928
## sex                 sex 0.10894407
## pornlaw2            pornlaw2 0.08501498
## black               black 0.00000000
## born               born 0.00000000
## degree_Bachelor.deg  degree_Bachelor.deg 0.00000000
## news_FEW.TIMES.A.WEEK news_FEW.TIMES.A.WEEK 0.00000000
## news_LESS.THAN.ONCE.WK news_LESS.THAN.ONCE.WK 0.00000000
## news_NEVER          news_NEVER 0.00000000
## news_other          news_other 0.00000000
## relig_CATHOLIC      relig_CATHOLIC 0.00000000
## relig_NONE          relig_NONE 0.00000000
## social_cons3_Mod    social_cons3_Mod 0.00000000
## zodiac_other        zodiac_other 0.00000000
```

Question 5 Apply to Test

As can be seen, our best model did a mediocre job on the test set, yet is not much a decrease compared to the training set. This means the overfitting problem is handled at least partially, yet the overall prediction accuracy could use more work.

```
boost_test_pred <- as.numeric(predict(boost,newdata = test))
error_rate <- sum(boost_test_pred - 1 == test$colrac)/nrow(test)
auc_test <- roc(test$colrac, boost_test_pred)$auc[1]
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
paste('correct classification rate is', error_rate)
```

```
## [1] "correct classification rate is 0.799188640973631"
```

```
paste('area under curve is', auc_test)
```

```
## [1] "area under curve is 0.792999006951341"
```

Question 6 Interpretation with ICE curve


For tolerance:

```
train = read.csv('gss_train.csv')
train$colrac = as.numeric(train$colrac) - 1
#levels(train$colrac) <- c('1','2') # Failed attempt to solve duplicate level problem

# Reloading train to reverse the type change to $colrac that is causing errors
# by 'duplicated levels'. The code below treat our colrac as numeric.

ice1 <- ice(
  object = boost$finalModel,
  X = train,
  y = train$colrac,
  predictor = which(colnames(train) == 'tolerance'),
  verbose = FALSE,
  #logodds = TRUE, This should be included for classification,
  # but I'm not doing it for the prolem metioned above.
  n.trees = 150
  # An arbitrary large value because the number of trees
  # used cannot be found at boost$finalmodel descriptions
)


plot(ice1,
     xlab = 'Tolerance',
     ylab = 'partial yhat (colrac)')
```



As can be seen from the plot, the level of tolerance has a positive effect on predicted outcome of whether allowing racist to teach at school. As the tolerance level increases, the yhat also increases unstably until tolerance level reaches 8. After 8, the effect of tolerance increase does not seem to have any further impact on the dependant variable.

Below is the same plot with shallower ylim to further illustrate the conditional effect.


```
plot(ice1,
      xlab = 'Tolerance',
      ylab = 'partial yhat (colrac)',
      ylim = c(0,1)
)
```

For age:

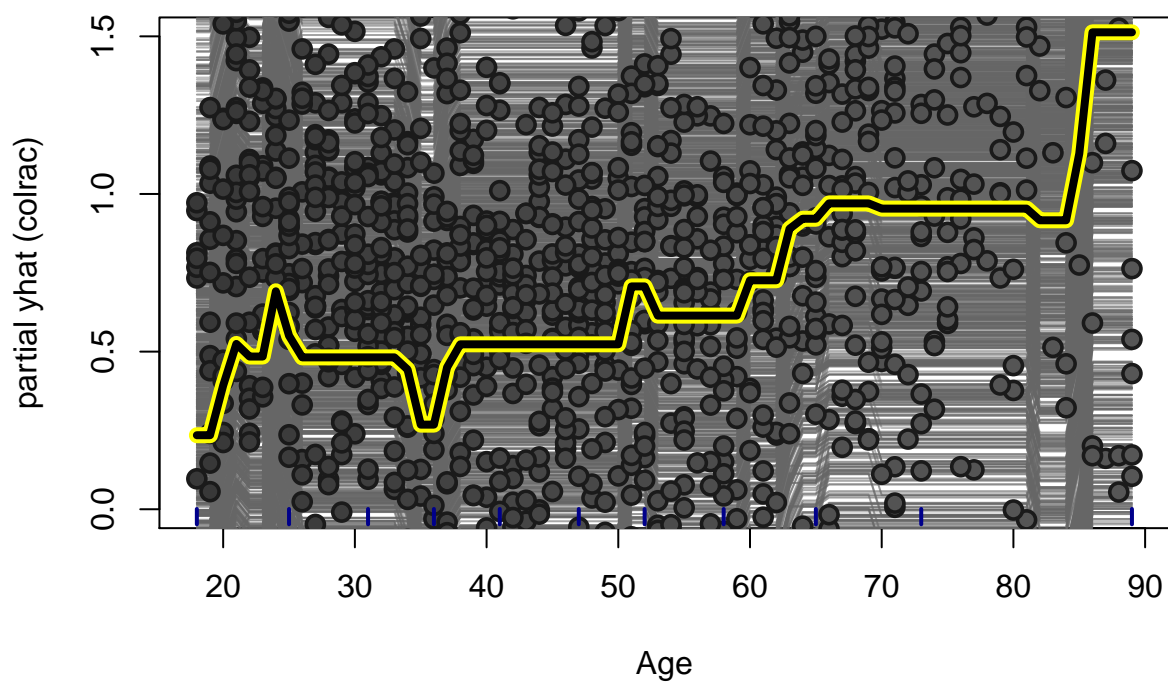
```
ice2 <- ice(
  object = boost$finalModel,
  X = train,
  y = train$colrac,
  predictor = 'age',
  verbose = FALSE,
  #logodds = TRUE, This should be included for classification,
  # but I'm not doing it for the prolem metioned above.
  n.trees = 150)

plot(ice2,
  xlab = 'Age',
  ylab = 'partial yhat (colrac)')
```



Zoomed-in version:

```
plot(ice2,  
     xlab = 'Age',  
     ylab = 'partial yhat (colrac)',  
     ylim = c(0,1.5)  
)
```



As can be seen from the plot, age on average has a positive effect on whether let racists teachers stay. The effect is rather unstable and weak at before 60 years old and we've seen a drop on mid 30s (where most parents have their children at school age). After 60 years old the effect becomes strong and stable and there's a spike at 80+.